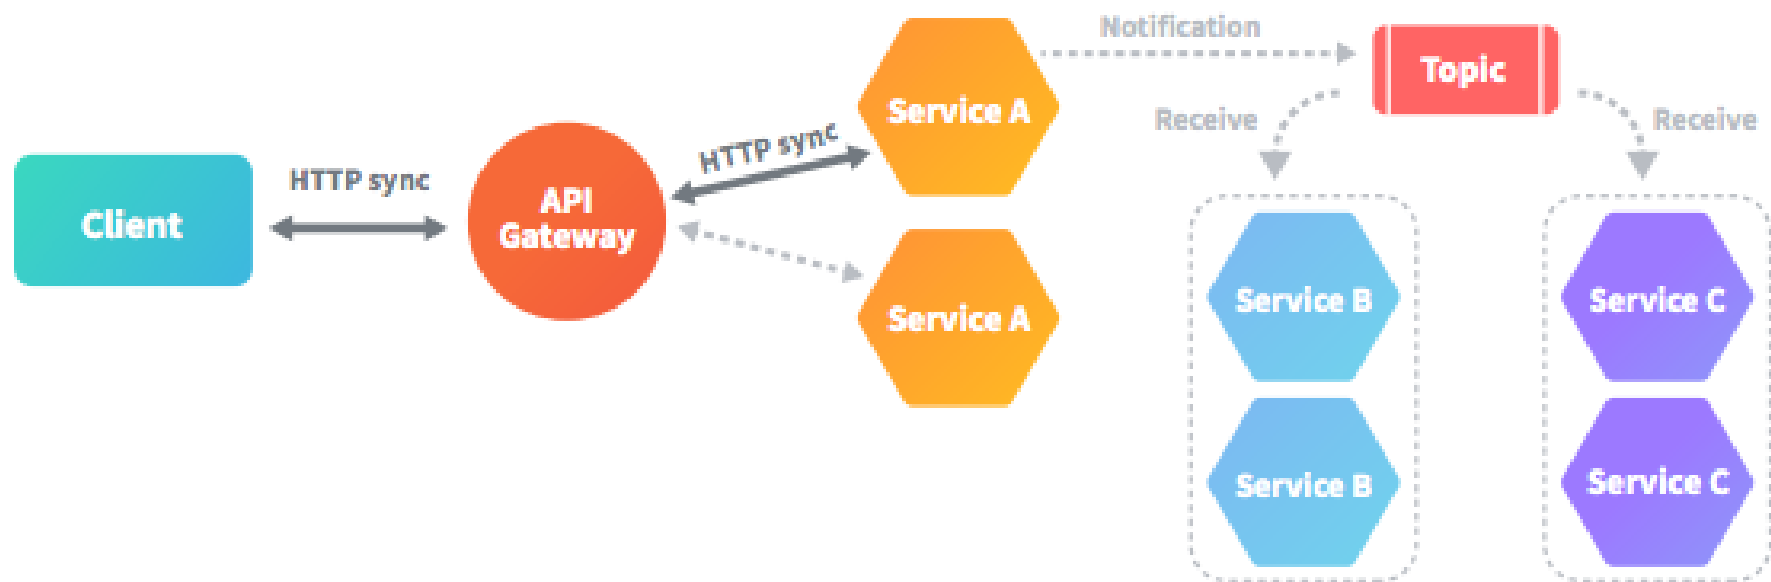


Apache Kafka

Bùi Minh Hoàng Vũ
Trần La Nhật Phương

Microservices

- Applications are built as loosely coupled services
- The services are heterogeneous in nature
- To work together, they need a way to communicate
 - Synchronous One-To-One
 - Asynchronous
 - Callback
 - Polling
 - Message Queue



Apache Kafka

- A distributed streaming platform
- Data within Kafka is stored:
 - Durably
 - In order
 - Can be read deterministically
- Can be distributed between servers

Publish-Subscribe Model

- The sender does not direct messages to a receiver
- The message is classified instead
- The receiver subscribes to classes of messages
- ➔ Decoupled Communication
- ➔ Dynamic Scaling
- ➔ Fault Tolerance

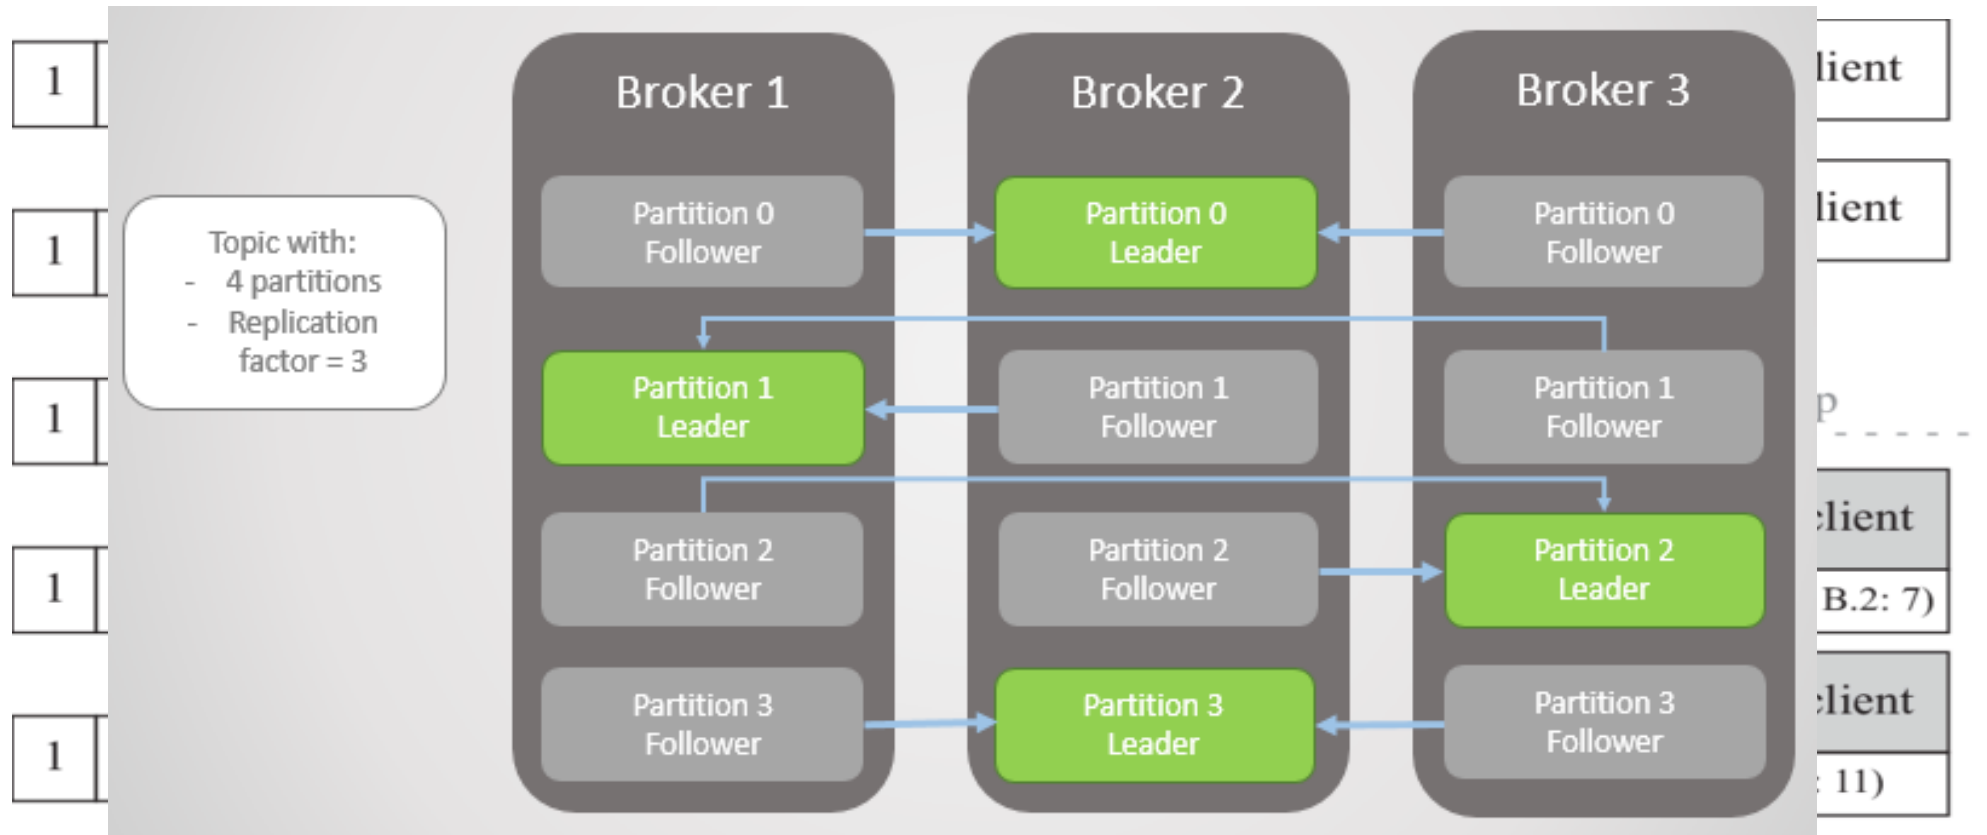
Core Concepts

- Message: A unit of data. Can optionally have a key.
- Schema: JSON, XML, Avro,...
- Topics: Organize events into categories
- Partitions: Only order within a partition is guaranteed.
- Brokers: Store and manage events.
- Producers and Consumers

Data Flow

- A producer publish a message to a specific topic
- Kafka Brokers store and replicate the message
 - Messages are assigned to partitions using the key
 - Messages are appended to the end of partitions
- Consumers read messages from specific partitions
- Messages are not deleted upon read

Partitions & Replication

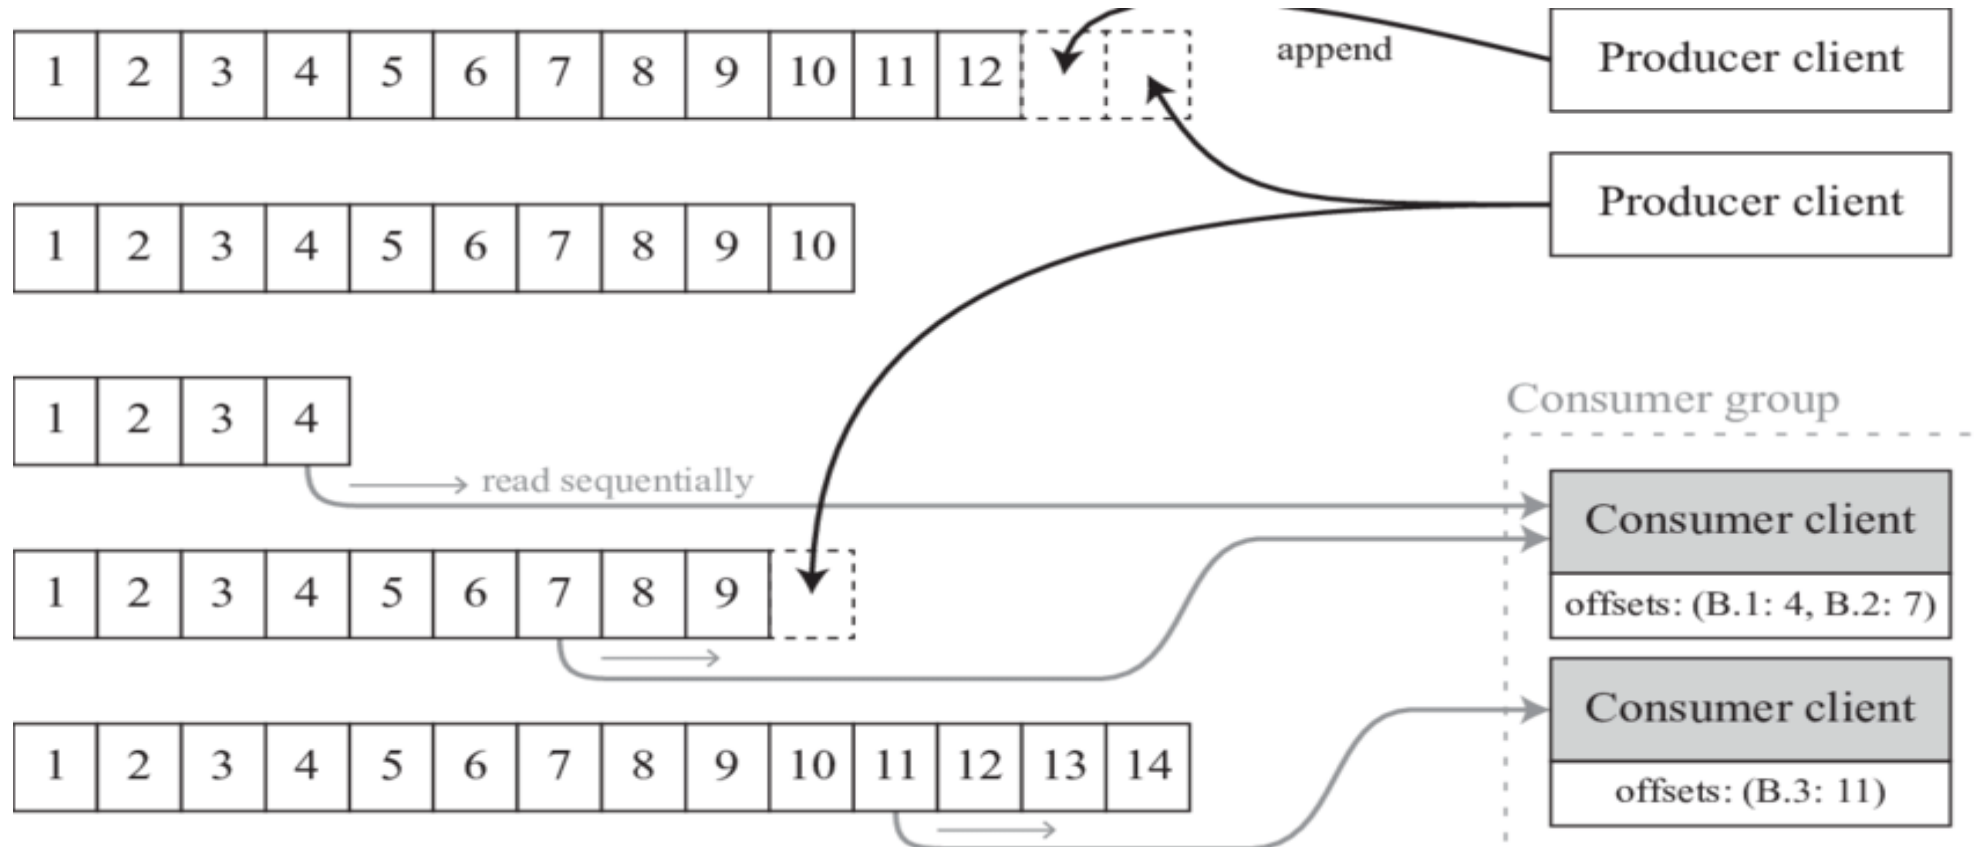


Producers

- Create & send messages to Topics
- Can use keys to influence partitioning
- Message acknowledgment:
 - No Ack
 - Leader Ack
 - All Replicas Ack

Consumers

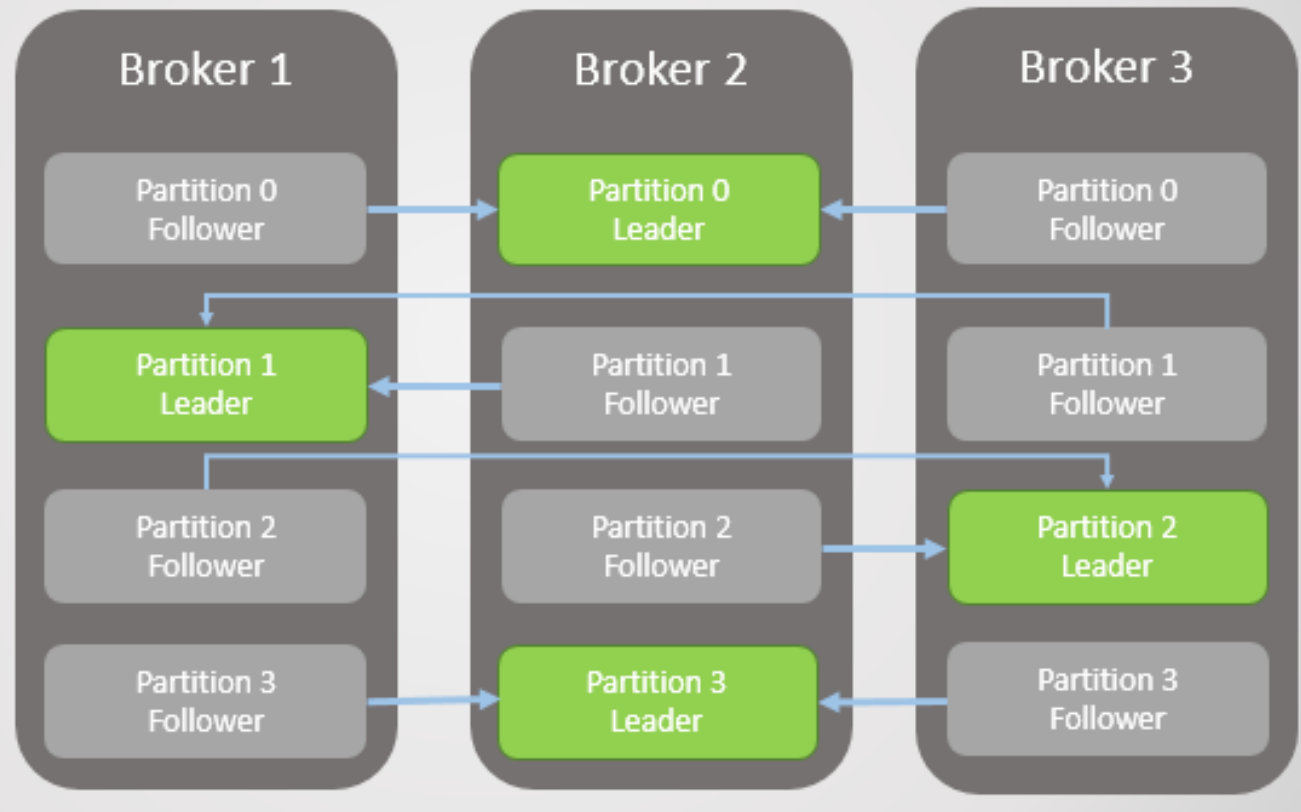
- Subscribe to topics and process messages
- Can be organized into groups
 - Each partition in a topic is consumed by only one consumer within the group
- Consumers maintain an offset
- Dynamic addition or removal of consumers



Brokers and Clusters

- A single Kafka server is called a Broker
- Brokers are designed to operate as part of a Cluster
 - One broker will function as the cluster's controller
 - Each partition is owned by a single cluster
- Retention: Messages only expire after some time (7 days) or the partition reach a certain size (1GB)

Topic with:
- 4 partitions
- Replication factor = 3



Use cases

- Activity Tracking
- Messaging
- Metrics and Logging
- Commit Logging
- Stream Processing

Pros & Cons

- ✓ Multiple producers, multiple consumers
- ✓ Disk-based retention
- ✓ Scalable
- ✓ High performance
- ✗ Learning curve
- ✗ Resource overhead
- ✗ Operational Complexity
- ✗ Overkill for small data

The Kafka Ecosystem

- Kafka Connect
- Kafka Streams
- Schema Registry
- MirrorMaker
- ZooKeeper – KRaft Mode

Thank you!